

LTLf Synthesis Under Environment Specifications¹

Antonio Di Stasio

Sapienza University of Rome

ICTCS 2022



¹Joint work with Giuseppe De Giacomo, Lucas Tabajara, Moshe Vardi and Shufang Zhu

Given a specification φ over inputs I and outputs O , expressed in:

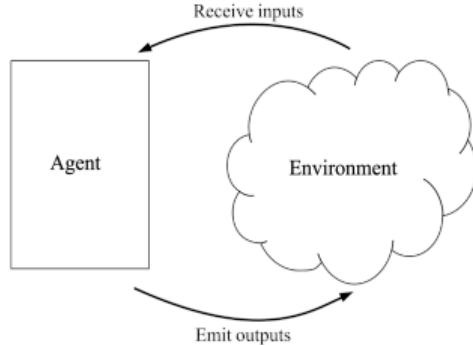
LTL (Pnueli 1977) or LTL_f (De Giacomo, Vardi 2013)

Syntax:

$$\varphi ::= a \mid \varphi \wedge \varphi \mid \neg \varphi \mid \circ \varphi \mid \varphi U \varphi \mid \Diamond \varphi \mid \Box \varphi$$

Semantic:

A trace $trace$ is an infinite (LTL) or finite (LTL_f) sequence over I and O . We write $trace \models \varphi$ to mean that τ satisfies φ .



Agent and Environment Strategies, and Traces

For an agent strategy $\sigma_{ag} : I^+ \rightarrow O$ and an environment strategy $\sigma_{env} : O^* \rightarrow I$

$$\text{trace}(\sigma_{ag}, \sigma_{env}) = (i_1 \cup o_1), (i_2 \cup o_2) \dots \in 2^{I \cup O}$$

denotes the unique trace induced by both σ_{ag} and σ_{env} .

Synthesis Problem

Given an LTL/ LTL_f task *Goal* for the agent

Find agent strategy σ_{ag} such that $\forall \sigma_{env}. \text{trace}(\sigma_{ag}, \sigma_{env}) \models \text{Goal}$

Algorithm for LTL synthesis

Given LTL formula φ

1. Compute corresponding Nondeterministic Büchi Automaton (NBA) (exponential)
2. Determinize NBA into Deterministic Parity Automaton (DPA) (exp in states, poly in priorities (nested-fixpoints))
3. Synthesize winning strategy for Parity Game (poly in states, exp in priorities)

Complexity

LTL synthesis is **2EXPTIME**-complete

Tools

- **Spot**^a: a platform for LTL and ω -automata manipulation.
- **Strix**^b: So far, the best tool for solving LTL synthesis.

^a<https://spot.lrde.epita.fr/>^b<https://strix.model.in.tum.de/>

Algorithm for LTL_f synthesisGiven LTL_f formula φ

1. Compute corresponding Nondeterministic Finite Automaton (NFA) (exponential)
2. Determinize NFA to Deterministic Finite Automaton (DFA) (exponential)
3. Synthesize winning strategy for DFA game (linear)

Complexity

LTL_f synthesis is 2EXPTIME-complete

Tools

- **Itlf2dfa**^a: a tool for translating LTL_f into DFA.
- **Syft, Lysa, Lydia, Cynthia**^b...

^a<http://ltlf2dfa.diag.uniroma1.it>.^b<https://github.com/whitemech>

Domain

- ▶ Planning consider the agent acting in a **(nondeterministic) domain**
- ▶ The domain is a **model of how the world** (i.e. the environment) works
- ▶ That is, it is a **specification of the possible environment strategies**

$$[[Dom]] = \{\sigma_{env} | \sigma_{env} \text{ compliant with } Dom\}$$

Planning in nondeterministic domains

Given an task *Goal* for the agent, and a domain *Dom* modeling the environment

Find agent behavior σ_{ag} such that $\forall \sigma_{env} \in [[Dom]].trace(\sigma_{ag}\sigma_{env}) \models Goal$



Agent Tasks terminate: LTL_f

- ▶ Because it is the agent that is planning/reasoning.
- ▶ If the task would not terminate, the agent would be stuck into doing the same task forever.
- ▶ We want to focus on autonomous intelligent agents that (1) get a task, (2) reason/plan autonomously to solve it, (3) execute the plan, (4) get another task, and so on.

We cannot require the environment specification to be in LTL_f : we must use LTL

The environment has to react to the agent's moves anyway, independently of whether the agent accomplishes its task (in a finite number of steps or not)

Which kinds of environment assumptions can the agent make?

- ▶ Nondeterministic planning domains;
- ▶ Forms of fairness ($\square\lozenge\phi$) and stability ($\lozenge\square\phi$) [ZhuDeGiacomoPuVardiAAI2020];
- ▶ Safety properties [DeGiacomoDiStasioPerelliZhuKR2021];
- ▶ GR(1) formulas [DeGiacomoDiStasioTabajaraVardiZhuIJCAI2021];
- ▶ ..

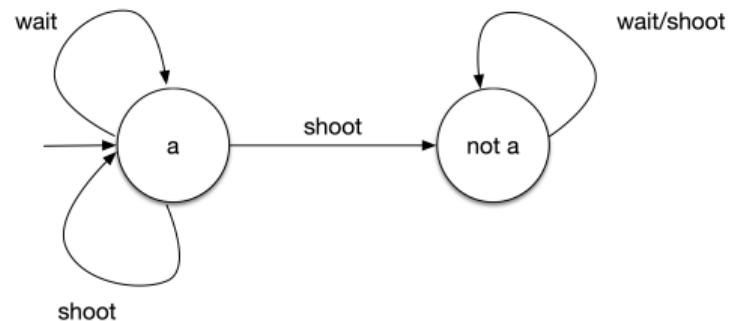
Environments Specifications as LTL formulas

A natural generalization is to consider general environment specifications expressed as arbitrary LTL formulas [DeGiacomoDiStasioVardiZhuKR2020].

For example let the environment specification be formed by $Env_1 \wedge Env_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

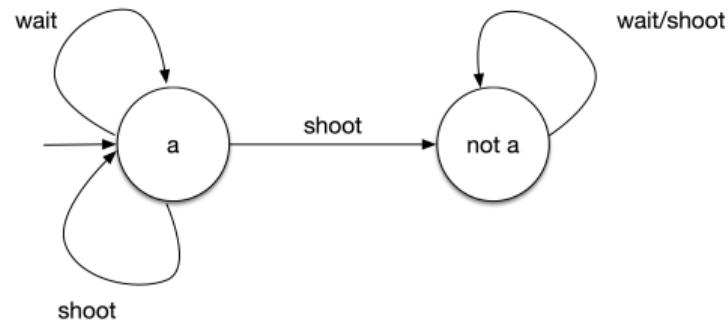
- $\square(alive \rightarrow \circ(wait \rightarrow alive))$
- $\square(alive \rightarrow \circ(shoot \rightarrow (alive \vee \neg alive)))$
- $\square(\neg alive \rightarrow \circ(wait \rightarrow \neg alive))$
- $\square(\neg alive \rightarrow \circ(shoot \rightarrow \neg alive))$
- $\square((wait \wedge shoot) \wedge (wait \rightarrow \neg shoot))$



For example let the environment specification be formed by $Env_1 \wedge Env_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

- $\square(alive \rightarrow \circ(wait \rightarrow alive))$
- $\square(alive \rightarrow \circ(shoot \rightarrow (alive \vee \neg alive)))$
- $\square(\neg alive \rightarrow \circ(wait \rightarrow \neg alive))$
- $\square(\neg alive \rightarrow \circ(shoot \rightarrow \neg alive))$
- $\square((wait \wedge shoot) \wedge (wait \rightarrow \neg shoot))$



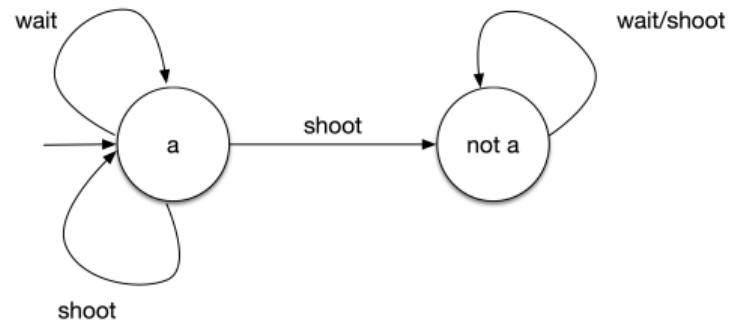
Env_2 is the LTL formula expressing some fairness over the nondeterministic effect a , e.g.,

$$\square \diamondsuit shoot \rightarrow \diamondsuit \neg a$$

For example let the environment specification be formed by $Env_1 \wedge Env_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

- $\square(alive \rightarrow \circ(wait \rightarrow alive))$
- $\square(alive \rightarrow \circ(shoot \rightarrow (alive \vee \neg alive)))$
- $\square(\neg alive \rightarrow \circ(wait \rightarrow \neg alive))$
- $\square(\neg alive \rightarrow \circ(shoot \rightarrow \neg alive))$
- $\square((wait \wedge shoot) \wedge (wait \rightarrow \neg shoot))$



Env_2 is the LTL formula expressing some fairness over the nondeterministic effect a , e.g.,

$$\square \lozenge shoot \rightarrow \lozenge \neg a$$

Let $Goal$ be an LTL_f formula which expresses an agent task, e.g.,

$$\lozenge \neg a$$

Environment specifications in LTL/LTL_f

Let Env be an LTL/LTL_f formula over \mathcal{I} and \mathcal{O} .

$$[[\text{Env}]] = \{\sigma_{\text{env}} \mid \forall \sigma_{\text{ag}}. \text{trace}(\sigma_{\text{ag}}, \sigma_{\text{env}}) \models \text{Env}\}$$

i.e Env denotes all environment strategies that play according to the specification whatever is the agent strategy.

Synthesis under environment specifications in LTL/LTL_f

Given an LTL_f task Goal for the agent, and an LTL/LTL_f environment specification Env :

Find agent strategy σ_{ag} such that $\forall \sigma_{\text{env}} \in [[\text{Env}]]. \text{trace}(\sigma_{\text{ag}}, \sigma_{\text{env}}) \models \text{Goal}$

Consistent environment specifications

Is any LTL/LTL_f formula a valid environment specification? No, *Env* needs to be "consistent"!:

$$[[Env]] \neq \emptyset$$

i.e. $\exists \sigma_e. \forall \sigma_{ag}. trace(\sigma_{ag}, \sigma_e) \models Env$



Environment Specifications

Let Env be an LTL/LTL_f formula over $I \cup O$.

$$[[\text{Env}]] = \{\sigma_{\text{env}} | \sigma_{\text{env}} \text{ satisfies } \text{Env} \text{ whatever is the agent strategy}\}$$

Synthesis under environment specifications in LTL/LTL_f

Given an LTL/ LTL_f task Goal for the agent, and an LTL/LTL_f environment specification Env :

Find agent strategy σ_{ag} such that $\forall \sigma_{\text{env}} \in [[\text{Env}]]. \text{trace}(\sigma_{\text{ag}}, \sigma_{\text{env}}) \models \text{Goal}$

Theorem [AminofDeGiacomoMuranoRubinICAPS2019]

To find agent strategy realizing Goal under the environment specification Env , we can use standard LTL/LTL_f synthesis for

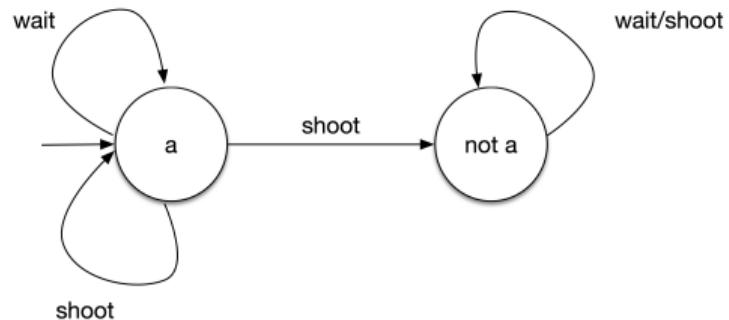
$$\text{Env} \rightarrow \text{Goal}$$

LTL_f Synthesis Under LTL Environment Specifications

For example let the assumption be formed by $\text{Env}_1 \wedge \text{Env}_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

- $\square(\text{alive} \rightarrow \circ(\text{wait} \rightarrow \text{alive}))$
- $\square(\text{alive} \rightarrow \circ(\text{shoot} \rightarrow (\text{alive} \vee \neg\text{alive}))$
- $\square(\neg\text{alive} \rightarrow \circ(\text{wait} \rightarrow \neg\text{alive}))$
- $\square(\neg\text{alive} \rightarrow \circ(\text{shoot} \rightarrow \neg\text{alive}))$
- $\square((\text{wait} \wedge \text{shoot}) \wedge (\text{wait} \rightarrow \neg\text{shoot}))$



Env_2 is the LTL formula expressing some fairness over nondeterministic effects, e.g.,

$$\square \lozenge \text{shoot} \rightarrow \lozenge \neg a$$

Let Goal be an LTL_f formula which expresses an agent task, e.g.,

$$\lozenge \neg a$$

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

Problem

Solve the synthesis problem for

$$\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$.

... but we can exploit the simplicity of dealing with LTL_f given:

Problem

Solve the synthesis problem for

$$\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$.

... but we can exploit the simplicity of dealing with LTL_f given:

- Env_1 : LTL

Problem

Solve the synthesis problem for

$$\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$.

... but we can exploit the simplicity of dealing with LTL_f given:

- Env_1 : LTL → LTL_f

Problem

Solve the synthesis problem for

$$\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$.

... but we can exploit the simplicity of dealing with LTL_f given:

- Env_1 : LTL → LTL_f
- Env_2 : LTL

Problem

Solve the synthesis problem for

$$\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $\text{Env}_1 \wedge \text{Env}_2 \rightarrow \text{Goal}$.

... but we can exploit the simplicity of dealing with LTL_f given:

- Env_1 : LTL \rightarrow LTL_f
- Env_2 : LTL
- Goal : LTL_f

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Problem

Solve the synthesis problem for

$$Env_2 \rightarrow Goal'$$

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Problem

Solve the synthesis problem for

$$Env_2 \rightarrow Goal'$$

How can we exploit that $Goal'$ is LTL_f?

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Problem

Solve the synthesis problem for

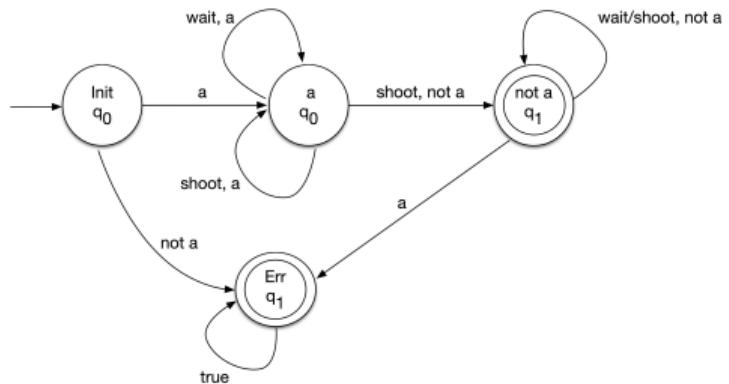
$$Env_2 \rightarrow Goal'$$

How can we exploit that $Goal'$ is LTL_f?

Two-stage technique!

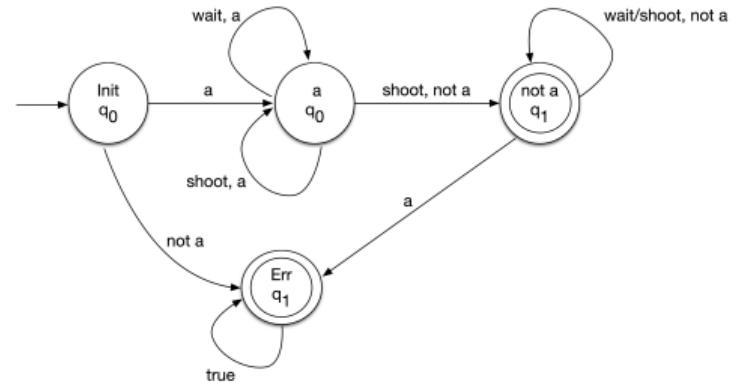
1° Stage

- ▶ Compute the corresponding DFA \mathcal{A} of $Goal' = \neg Env_1 \vee Goal$.



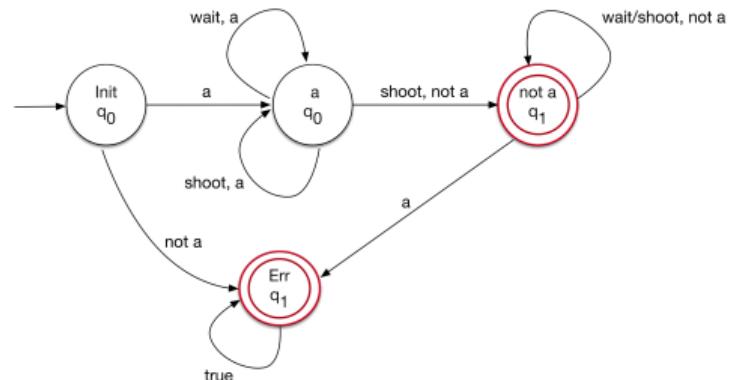
1° Stage

- ▶ Compute the corresponding DFA \mathcal{A} of $Goal' = \neg Env_1 \vee Goal$.
- ▶ Solve the reachability game for the agent over \mathcal{A} .



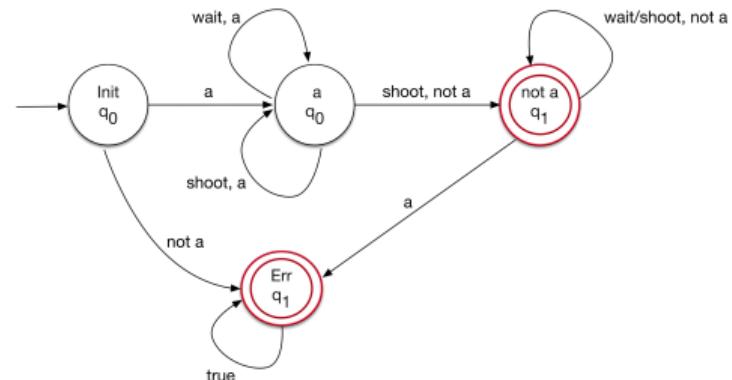
1° Stage

- ▶ Compute the corresponding DFA \mathcal{A} of $Goal' = \neg Env_1 \vee Goal$.
- ▶ Solve the reachability game for the agent over \mathcal{A} .
- ▶ Check whether the initial state is winning for the agent.



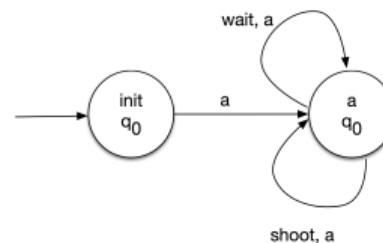
1° Stage

- ▶ Compute the corresponding DFA \mathcal{A} of $Goal' = \neg Env_1 \vee Goal$.
- ▶ Solve the reachability game for the agent over \mathcal{A} .
- ▶ Check whether the initial state is winning for the agent.
- ▶ If the initial state is not winning go to Stage 2, otherwise return the agent winning strategy.



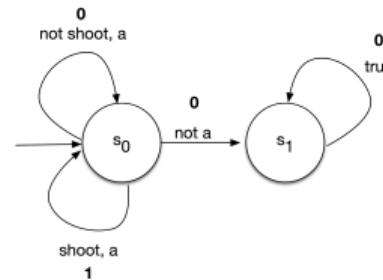
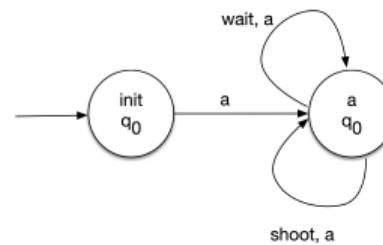
2° Stage

- ▶ Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .



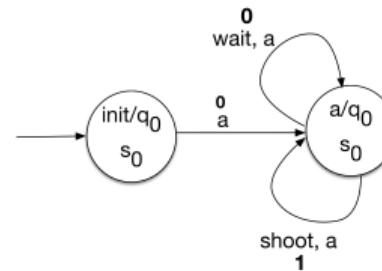
2° Stage

- ▶ Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
- ▶ Compute the corresponding DPA \mathcal{B} of Env_2 .



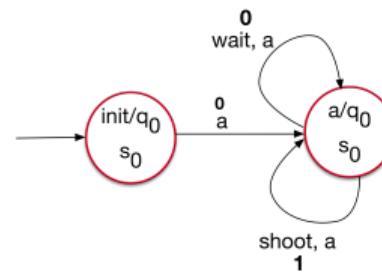
2° Stage

- ▶ Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
- ▶ Compute the corresponding DPA \mathcal{B} of Env_2 .
- ▶ Do the cartesian product between \mathcal{A}' and \mathcal{B} .



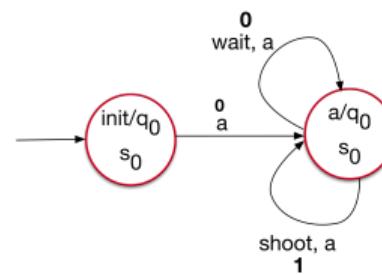
2° Stage

- ▶ Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
- ▶ Compute the corresponding DPA \mathcal{B} of Env_2 .
- ▶ Do the cartesian product between \mathcal{A}' and \mathcal{B} .
- ▶ Solve the parity game for the environment over $\mathcal{A}' \times \mathcal{B}$.



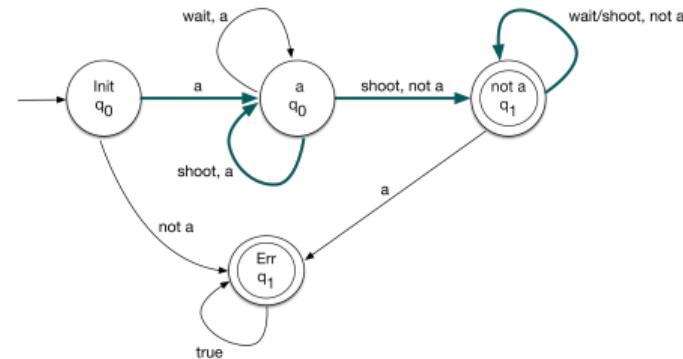
2° Stage

- ▶ Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
- ▶ Compute the corresponding DPA \mathcal{B} of Env_2 .
- ▶ Do the cartesian product between \mathcal{A}' and \mathcal{B} .
- ▶ Solve the parity game for the environment over $\mathcal{A}' \times \mathcal{B}$.
- ▶ Check if the initial state is winning for the agent; if not return "Unrealizable".



2° Stage

- ▶ Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
- ▶ Compute the corresponding DPA \mathcal{B} of Env_2 .
- ▶ Do the cartesian product between \mathcal{A}' and \mathcal{B} .
- ▶ Solve the parity game for the environment over $\mathcal{A}' \times \mathcal{B}$.
- ▶ Check if the initial state is winning for the agent; if not return "Unrealizable".
- ▶ Return the agent winning strategy by combining the agent winning strategies in Stage 1 and 2.



LTL_f Synthesis Under GR(1) Environment Specifications

- ▶ GR(1) formula, generalization of fairness
 - ▶ Powerful expressiveness on environment assumption
- ▶ LTL_f formula, natural for finite-horizon task
 - ▶ Strong agent task specification
- ▶ Simple solution
 - ▶ No detour to LTL synthesis
 - ▶ Games on finite-word automata

- ▶ GR(1) game $\mathcal{G} = (\mathcal{A}, \psi)$
 - ▶ Game arena \mathcal{A} encoding moves
 - ▶ Winning condition GR(1) formula ψ , powerful notion of fairness

$$\psi = \bigwedge_{i=1}^m \square \diamond \mathcal{J}_i \rightarrow \bigwedge_{j=1}^n \square \diamond \mathcal{K}_j$$

- ▶ **Complexity:** simple game solving, quadratic time

Problem

Solve the synthesis problem for

$$\varphi_{GR(1)}^e \rightarrow \varphi_{task}^a$$

Key Idea

- ▶ Agent goal: $\neg\varphi_{GR(1)}^e \vee \varphi_{task}^a$ — Environment goal $\varphi_{GR(1)}^e \wedge \neg\varphi_{task}^a$.
- ▶ Build the corresponding DFA $\mathcal{A}_{\varphi_{task}^a}$ of φ_{task}^a , and take its complement $\overline{\mathcal{A}_{\varphi_{task}^a}}$.
- ▶ Define a GR(1) game whose game arena is $\overline{\mathcal{A}_{\varphi_{task}^a}}$ and winning condition $\varphi_{GR(1)}^e$.
- ▶ Solve the GR(1) game for the agent, i.e., solve the dual of the GR(1) game.

LTL_f Synthesis Under Safety Environment Specifications



Definition

A safety property is a property which specifies that some (bad) behavior will never occur.

Examples:

- ▶ "always at most one process is in its critical section"
- ▶ "money can only be withdrawn once a correct PIN has been provided"

Important property

Any infinite trace violating the property has a finite prefix that is "bad";

- ▶ ... two processes are in the critical section ...

Usually: $\Box\neg\ldots$

Nondeterministic Strategy

A *non-deterministic* strategy for the agent is defined as a function $\Pi : (2^{\mathcal{X}})^* \rightarrow 2^{\mathcal{Y}}$. The set of deterministic strategies induced by a non-deterministic strategy Π is the set $[[\Pi]] = \{\sigma : (2^{\mathcal{X}})^* \rightarrow 2^{\mathcal{Y}} \mid \sigma(h) \in \Pi(h), \text{ for } h \in (2^{\mathcal{X}})^*\}$.

Maximally Permissive Strategy

A non-deterministic strategy Π is at least as permissive as Π' if $[[\Pi']] \subseteq [[\Pi]]$. A non-deterministic strategy Π is a maximally permissive strategy if $[[\Pi']] \subseteq [[\Pi]]$, for every non-deterministic strategy Π' .

Theorem

Safety specifications admit a nondeterministic strategy Π that captures the maximally permissive strategy.²

²Bernet et al.: Permissive strategies: from parity games to safety games

Problem

Solve the synthesis problem for

$$\varphi_S^e \rightarrow \varphi_{task}^a$$

Key Idea

1. Compute the *deterministic safety automaton* \mathcal{D} of φ_S (**no Büchi determinization!**)
2. Solve the safety game for the **environment** over \mathcal{D} ;
3. Construct the *maximally permissive strategy* \mathcal{T} (transducer);
4. Build the corresponding DFA $\mathcal{A}_{\varphi_{task}^a}$ of φ_{task}^a ;
5. Do the product of \mathcal{T} and $\mathcal{A}_{\varphi_{task}^a}$;
6. Solve the reachability game for the **agent** over it, and return a strategy, if exists.

Almost all the techniques are based on the following reduction

Theorem [Aminof et al. ICAPS 2019]

To find agent strategy realizing *Goal* under the environment specification *Env*, we can use standard synthesis for

$$\textit{Env} \rightarrow \textit{Goal}$$

Possible directions

- ▶ In case of safety environment specifications we can directly solve the problem without reduction to the implication.
- ▶ What about the other environment specifications? Fairness, GR(1), LTL, ...